

Table of Contents

Overview	1
Using libyottadb	1
Concepts	2
Data Types	2
Symbolic Constants	2
<i>Function Return Codes</i>	2
Normal Return Codes	2
Error Return Codes	3
<i>Limits</i>	3
Data Structures	3
Macros	3
API	4
STUFF BELOW IS FROM OLD DOCUMENT	4
<i>Query</i>	4
<i>Update</i>	6
<i>Transaction Processing</i>	6
Programming Notes	6
<i>Dynamic typing with automatic conversion</i>	6
<i>Numeric Considerations</i>	6
Canonical Numbers	7
<i>Tokens</i>	8

Overview

This is user documentation for accessing the YottaDB engine from C using the Simple API. A process can both call the Simple API as well as call functions written in M and exported.

Caveat: This code does not exist yet. The user documentation is being written ahead of the code, and will change in the event the code needs to differ from this documentation.

Using libyottadb

1. Install YottaDB.
2. Include the `yottadb.h` file in your C program and compile it.
3. Perform any database configuration and initialization needed (configuring global directories, creating database files, starting a Source Server process, etc.).

4. Run your program, ensuring either that `libyottadb.so` is in the load path of your program, or that it is preloaded.

Concepts

Key-value

Local and global variables

Subscripts (keys) of variables accessed using Simple API are strings. When a string is a [canonical number](#) YottaDB internally converts and stores it as a number. When ordering (collating) subscripts:

- Null (empty string) subscripts precede all numeric subscripts.
 - **YottaDB strongly recommends against applications that use null subscripts.**
- Numeric subscripts precede string subscripts.
 - Numeric subscripts in numeric order.
- String subscripts collate in byte order.

Data Types

Data types are defined by including `yottadb.h` and are one of:

`ydb_int_t` and `ydb_uint_t` — Signed and unsigned integers, that are *at least* 16 bits.

`ydb_long_t` and `ydb_ulong_t` — Signed and unsigned integers, that are *at least* 32 bits.

`ydb_longlong_t` and `ydb_ulonglong_t` — Signed and unsigned integers that are *at least* 64 bits. See [Numeric Considerations](#) below.

`ydb_status_t` — A signed integer which is the return value (status) of a call to a libyottadb function.

Symbolic Constants

The `yottadb.h` file defines several symbolic constants, which are one of the following types:

- Function Return Codes, which in turn are one of:
 - Normal Return Codes
 - Error Return Codes
- Limits
- Other

Function Return Codes

Return codes from calls to libyottadb are of type `ydb_status_t`. Normal return codes are non-negative (greater than or equal to zero); error return codes are negative.

Normal Return Codes

Symbolic constants for normal return codes are prefixed with `YDB_`.

`YDB_STATUS_OK` — Normal return following successful execution.

Error Return Codes

Symbolic constants for error codes returned by calls to libyottadb are prefixed with `YDB_ERR_`.¹ The symbolic constants below are not intended to be a complete list of all error messages that Simple API functions can return - the `ydb_message()` functions provides a way to get detailed information about a error codes for those without symbolic constants.

`YDB_ERR_GVINVALID` — A global variable name is too long.²

`YDB_ERR_GVUNDEF` — No value exists at a requested global variable node.

`YDB_ERR_LVUNDEF` — No value exists at a requested local variable node.³

`YDB_ERR_INVSTRLEN` — A buffer provided by the caller is not long enough for the string to be returned, or the length of a string passed as a parameter exceeds `YDB_MAX_STR`.

`YDB_ERR_KEY2BIG` — The length of a global variable name and subscripts exceeds the limit configured for a database region.

`YDB_ERR_MAXNRSUBSCRIPTS` — The number of subscripts specified in the call exceeded `YDB_MAX_SUB`.

`YDB_ERR_UNKNOWN` — A call to `ydb_zmessage()` specified an invalid message code.

Limits

Symbolic constants for limits are prefixed with `YDB_MAX_`. Unless otherwise noted, symbolic constants are unsigned integers guaranteed to fit within the range of a `ydb_uint_t` type.

`YDB_MAX_IDENT` — The maximum space in bytes required to store a complete variable name, including the preceding caret for a global variable.

`YDB_MAX_MSG` — The maximum length in bytes of any message string associated with a message code. A buffer of length `YDB_MAX_MSG` bytes ensures that a call to `ydb_zmessage()` will not return a `YDB_ERR_INVSTRLEN` return code.

`YDB_MAX_STR` — The maximum length of a string (or blob) in bytes. A caller to `ydb_get()` that provides a buffer of `YDB_MAX_STR` will never get a `YDB_ERR_INVSTRLEN` error. `YDB_MAX_STR` is guaranteed to fit in a `ydb_ulong_t` type.

`YDB_MAX_SUB` — The maximum number of subscripts for a local or global variable.

Data Structures

`ydb_string_t` is a descriptor for a string⁴ value, and consists of the following fields:

- `alloc` and `used` — fields of type `ydb_strlen_t` where `alloc` \geq `used`
- `address` — pointer to a `ydb_zchar_t`, the starting address of a string

Macros

`YDB_ALLOC_STRING(x, strlit)` — Allocate a `ydb_string_t` structure and initialize it to `strlit`, returning the address of the structure. Note that while the `used` field is the size of `strlit`, the `alloc` field may be rounded up to a larger value.⁵

`YDB_FREE_STRING(x)` — Free the `ydb_string_t` structure pointed to by `x`.

API

As YottaDB local and global variables can have variable numbers of subscripts, to allow the libyottadb Simple API functions to have variable numbers of parameters, the last parameter must always be `NULL` (the standard C symbolic constant).

```
ydb_status_t ydb_data(  
    ydb_uint_t *value,  
    ydb_string_t *varname,  
    [ ydb_string_t *subscript, ... ]  
    NULL);
```

In the location pointed to by `value`, returns the following information about the local or global variable node identified by `glvn`:

- 0 — There is neither a value nor a sub-tree, i.e., it is undefined.
- 1 — There is a value, but no sub-tree
- 10 — There is no value, but there is a sub-tree.
- 11 — There are both a value and a subtree.

STUFF BELOW IS FROM OLD DOCUMENT

Query

`ydb_status_t ydb_alias_handle(ydb_string_t *value, ydb_varsub_t *lvn)`

In the location pointed to by `value->address` returns the handle of the local variable referenced by `lvsub`. It is not meaningful for a caller to perform any operations on handles except to compare two handles for equality.

`ydb_status_t ydb_data(ydb_uint_t *value, ydb_varsub_t *glvn)`

In the location pointed to by `value`, returns the following information about the local or global variable node identified by `glvn`:

- 0 — There is neither a value nor a sub-tree, i.e., it is undefined.
- 1 — There is a value, but no sub-tree
- 10 — There is no value, but there is a sub-tree.
- 11 — There are both a value and a subtree.

The following values are only meaningful if `glvn` identifies a local variable node:

- 100 — The node is an alias, but there is neither a value nor a sub-tree.
- 101 — The node is an alias with a value but sub-tree.
- 110 — The node is an alias with no value, but with a sub-tree.
- 111 — The node is an alias with a value and a sub-tree.

`ydb_get(ydb_value_t *value, ydb_varsub_t *glivn)`

In the container pointer to by `value`, libyottadb returns the value referred to by `glivn`. If there is no value at the requested global or local variable node, or if the intrinsic special variable does not exist, a non-zero return value of `YDB_ERR_GVUNDEF`, `YDB_ERR_INVSVN`, or `YDB_ERR_UNDEF` indicates the error.

In a database application, a global variable node can potentially be changed by another process between the time that a process calls `ydb_length()` to get the length of the data in a node and a `ydb_get()` call to get that data. If a caller cannot ensure from the application schema that the size of the buffer it provides is large enough for a string returned by `ydb_get()`, it should code in anticipation of a potential `YDB_ERR_INVSTRLEN` return code from `ydb_get()`.

`ydb_length(ydb_ulong_t *value1, ydb_ulong_t *value2, ydb_ulong_t *value3, ydb_varsub_t *glivn)`

For each non-null `value*` parameter, in the memory location pointed to by that parameter, libyottadb returns the following information about the node specified by `glivsub`. If there is no value at the requested global or local variable node, or if a requested intrinsic special variable does not exist, a non-zero return value of `YDB_ERR_GVUNDEF`, `YDB_ERR_INVSVN`, or `YDB_ERR_UNDEF` indicates the error.

- `*value1` — the length of the data in bytes. If the data is numeric, `*value1` has the length of the canonical string representation of that value.
- `*value2` — if the environment variable `ydb_chset` at process startup has the (case insensitive) value “UTF-8”, `*value2` has the length of `*glivsub` in bytes; otherwise `*value2` is the same as `*value1`.
- `*value3` — if the environment variable `ydb_chset` at process startup has the (case insensitive) value “UTF-8”, `*value3` has the length of `*glivsub` in glyphs; otherwise `*value3` is the same as `*value1`.

`ydb_node_next(ydb_varsub_t *next, ydb_value_t *value, ydb_varsub_t *glvn)`

`ydb_node_next()` returns the next node in the tree in depth first search order, if one exists:

- If `next->varname->name->alloc ≥ glvn->varname->name->used` `ydb_node_next()` copies the `name->address` and `name->used` sub-fields from the `glvn->varname` structure to the

next->varname structure, returning a YDB_ERR_INVSTRLEN error if next->varname->name->alloc < glvn->varname->name->used.

- If next->varsub_alloc is large enough to hold the subscripts ydb_node_next() sets next->varsub_used to the actual number of subscripts, and

Update

Transaction Processing

Programming Notes

Dynamic typing with automatic conversion

The YottaDB engine internally automatically converts values between numbers and strings as needed. Thus it is legitimate to lexically compare the numbers 2 and 11, with the expected result that 11 precedes 2, and it is equally legitimate to numerically compare the strings “2” and “11”, with the expected result that 11 is greater than 2 — the functions for numeric and lexical comparisons are different.

In the ydb_value_t structure, a caller specifies what conversion, if any, it wishes the called libyottadb function to perform on the return value:

- When a value is numeric, and the requested type is a string (the tag field is YDB_STRING_STAR), libyottadb returns the number as a canonical string in the ydb_string_t structure pointed to by string_star
- When the value is a string, and the requested type is numeric, libyottadb converts

When returning a string, libyottadb functions *always* check that the alloc field of the ydb_string_t structure is large enough for the result, returning a YDB_ERR_STRLEN error if it is not.

Numeric Considerations

To ensure the accuracy of financial calculations, YottaDB internally stores nnumbers as, and performs arithmetic using, a scaled packed decimal representation with 18 significant decimal digits, with optimizations for values within a certain subset of its full range. Consequently:

- Any number that is exactly represented in YottaDB can be exactly represented as a string, with reasonably efficient conversion back and forth.
- Any integer value of up to 18 significant digits can be exactly represented by an integer type such as ydb_longlong_t, and integers in the inclusive range $\pm 999,999$ are handled more efficiently than larger integers.
- In YottaDB there are numbers which can be exactly represented (such as 0.1), but which cannot be exactly represented in binary floating point.
- In 64 bit integers and binary floating point formats, there are numbers which can be exactly represented, but which cannot be exactly represented in YottaDB.

This means that for numeric keys which are not guaranteed to be integers:

- In theory, there are edge cases where a value (which would internally be in YottaDB format) returned by a function such as `ydb_subscript_next()` and converted to a `ydb_double_t` when passed back to C application code, and then converted back to YottaDB internal format in a call to `ydb_get()` can result in the node not being found because the double conversion produces a number not identical to the original. Furthermore, there is a cost to the conversion.
- Passing keys back and forth as strings avoids those edge cases, but of course still has a conversion cost.

To preserve accuracy of numeric values that are returned by libyottadb, and which an application code intends to simply pass back to libyottadb as a libyottadb provides a `ydb_numeric_t` type. A value obtained from libyottadb in `ydb_numeric_t` loses no precision when returned to libyottadb, and as a further benefit is very efficient. While the actual value of `ydb_numeric_t` is opaque to application code, the `ydb_convert()` function is available.

Conversely, when passed a string that is a **canonical number** for use as a key, libyottadb automatically converts it to a number. This automatic internal conversion is irrelevant for the majority of typical application that:

- simply store and retrieve data associated with keys, potentially testing for the existence of nodes; or
- transfer keys which are numeric values between application code and libyottadb using numeric types and expect numeric ordering.

However, this automatic internal conversion does affect applications that:

- use numeric keys and expect the keys to be sorted in lexical order rather than numeric order; or
- transfer keys which are numeric values between application code and libyottadb as strings that may or may not be canonical numbers.

Applications that are affected by automatic internal conversion should prefix their keys with a character such as “x” which ensures that keys are not canonical numbers.

Canonical Numbers

Conceptually, a canonical number is a string from the Latin character set that represents a decimal number in a standard, concise, form.

1. Any string of decimal digits, optionally preceded by a minus sign (“-“), the first of which is not “0” (except for the number zero itself), that represents an integer of no more than 18 significant digits.
 - The following are canonical numbers: “-1”, “0”, “3”, “10”, “999999999999999999”, “9999999999999999990”. Note that the last string has only 18 significant digits even though it is 19 characters long.
 - The following are not canonical numbers: “+1” (starts with “+”), “00” (has an extra leading zero), “9999999999999999999” (19 significant digits).
2. Any string of decimal digits, optionally preceded by a minus sign that includes one decimal point (“.”), the first and last of which are not “0”, that represents a number of no more than 18 significant digits.
 - The following are canonical numbers: “-.1”, “.3”, “.999999999999999999”.

- The following are not canonical numbers “+.1” (starts with “+”), “0.3” (first digit is “0”), “.9999999999999999990” (last digit is “0”), “.9999999999999999999” (more than 18 significant digits).
3. Any of the above two forms followed by “E” followed by a canonical number integer in the range -43 to +47 such that the magnitude of the resulting number is between 1E-43 through 1E47.

Tokens

Since numeric and non-numeric subscripts can be freely intermixed in YottaDB, it requires knowledge of the application schema to know whether an application mixes numeric and string subscripts at the same level for a variable.

Consider whether this can be deferred for an initial implementation.

-
- | | |
|---|---|
| 1 | Note for implementers: the actual values are negated ZMESSAGE error codes. |
| 2 | Note for implementers: YottaDB silently truncates local variable names that are too long. The implementation should catch this and return an error code, e.g., something like YDB_ERR_LVINVALID. |
| 3 | Note for implementers: under the covers, this is UNDEF but renamed to be more meaningful. |
| 4 | Strings in YottaDB are arbitrary sequences of bytes that are not null-terminated. Other languages may refer to them as binary data or blobs. |
| 5 | Note for implementers: under the covers, YDB_ALLOC_STRING() and YDB_FREE_STRING() should call the ydb_malloc() and ydb_free() functions, which are aliases for the gtm_malloc() and gtm_free() functions (i.e., either prefix calls the same function). |