

# Table of Contents

<b>Overview</b>	<b>1</b>
<b>Using libyottadb</b>	<b>1</b>
<b>Concepts</b>	<b>2</b>
<b>Data Types</b>	<b>2</b>
<b>Symbolic Constants</b>	<b>2</b>
<i>Function Return Codes</i>	2
Normal Return Codes	3
Error Return Codes	3
<i>Limits</i>	3
<b>Data Structures</b>	<b>3</b>
<b>Macros</b>	<b>4</b>
<b>Simple API</b>	<b>4</b>
<b>STUFF BELOW IS FROM OLD DOCUMENT</b>	<b>5</b>
<i>Query</i>	5
<i>Update</i>	6
<i>Transaction Processing</i>	6
<b>Programming Notes</b>	<b>7</b>
<i>Dynamic typing with automatic conversion</i>	7
<i>Numeric Considerations</i>	7
Canonical Numbers	8
<i>Tokens</i>	8

## Overview

This is user documentation for accessing the YottaDB engine from C using the Simple API. A process can both call the Simple API as well as call functions written in M and exported.

**Caveat:** This code does not exist yet. The user documentation is being written ahead of the code, and will change in the event the code needs to differ from this documentation.

## Using libyottadb

1. Install YottaDB.
2. Include the `yottadb.h` file in your C program and compile it.
3. Perform any database configuration and initialization needed (configuring global directories, creating database files, starting a Source Server process, etc.).

4. Run your program, ensuring either that `libyottadb.so` is in the load path of your program, or that it is preloaded.

## Concepts

Key-value

Local and global variables

Subscripts (keys) of variables accessed using Simple API are strings. When a string is a [canonical number](#) YottaDB internally converts and stores it as a number. When ordering (collating) subscripts:

- Null (empty string) subscripts precede all numeric subscripts.
  - **YottaDB strongly recommends against applications that use null subscripts.**
- Numeric subscripts precede string subscripts.
  - Numeric subscripts in numeric order.
- String subscripts collate in byte order.

## Data Types

Data types are defined by including `yottadb.h` and are one of:

`ydb_int_t` and `ydb_uint_t` — Signed and unsigned integers, that are *at least* 16 bits.

`ydb_long_t` and `ydb_ulong_t` — Signed and unsigned integers, that are *at least* 32 bits.

`ydb_longlong_t` and `ydb_ulonglong_t` — Signed and unsigned integers that are *at least* 64 bits. See [Numeric Considerations](#) below.

`ydb_status_t` — A signed integer which is the return value (status) of a call to a libyottadb function.

`ydb_strlen_t` — An unsigned integer type that is able to store the maximum length of a string, `YDB_MAX_STR`.

`ydb_uchar_t` — An unsigned data value that is *exactly* 8-bits (one byte).

## Symbolic Constants

The `yottadb.h` file defines several symbolic constants, which are one of the following types:

- Function Return Codes, which in turn are one of:
  - Normal Return Codes
  - Error Return Codes
- Limits
- Other

### *Function Return Codes*

Return codes from calls to libyottadb are of type `ydb_status_t`. Normal return codes are non-negative (greater than or equal to zero); error return codes are negative.

## Normal Return Codes

Symbolic constants for normal return codes are prefixed with `YDB_`.

`YDB_STATUS_OK` — Normal return following successful execution.

## Error Return Codes

Symbolic constants for error codes returned by calls to libyottadb are prefixed with `YDB_ERR_`.<sup>1</sup> The symbolic constants below are not intended to be a complete list of all error messages that Simple API functions can return - the `ydb_message()` functions provides a way to get detailed information about a error codes for those without symbolic constants.

`YDB_ERR_GVINVALID` — A global variable name is too long.<sup>2</sup>

`YDB_ERR_GVUNDEF` — No value exists at a requested global variable node.

`YDB_ERR_LVUNDEF` — No value exists at a requested local variable node.<sup>3</sup>

`YDB_ERR_INVSTRLEN` — A buffer provided by the caller is not long enough for the string to be returned, or the length of a string passed as a parameter exceeds `YDB_MAX_STR`. In the event the return code is `YDB_ERR_INVSTRLEN` (for which the first parameter must be of the form `ydb_string_t *value`), then `value->used` is set to the size required of a sufficiently large buffer, and `value->address` points to the first `value->alloc` bytes of the value.

`YDB_ERR_KEY2BIG` — The length of a global variable name and subscripts exceeds the limit configured for a database region.

`YDB_ERR_MAXNRSUBSCRIPTS` — The number of subscripts specified in the call exceeded `YDB_MAX_SUB`.

`YDB_ERR_UNKNOWN` — A call to `ydb_zmessage()` specified an invalid message code.

## Limits

Symbolic constants for limits are prefixed with `YDB_MAX_`. Unless otherwise noted, symbolic constants are unsigned integers guaranteed to fit within the range of a `ydb_uint_t` type.

`YDB_MAX_IDENT` — The maximum space in bytes required to store a complete variable name, including the preceding caret for a global variable.

`YDB_MAX_MSG` — The maximum length in bytes of any message string associated with a message code. A buffer of length `YDB_MAX_MSG` bytes ensures that a call to `ydb_zmessage()` will not return a `YDB_ERR_INVSTRLEN` return code.

`YDB_MAX_STR` — The maximum length of a string (or blob) in bytes. A caller to `ydb_get()` that provides a buffer of `YDB_MAX_STR` will never get a `YDB_ERR_INVSTRLEN` error. `YDB_MAX_STR` is guaranteed to fit in a `ydb_ulong_t` type.

`YDB_MAX_SUB` — The maximum number of subscripts for a local or global variable.

## Data Structures

`ydb_string_t` is a descriptor for a string<sup>4</sup> value, and consists of the following fields:

- `alloc` and `used` — fields of type `ydb_strlen_t` where `alloc ≥ used`
- `address` — pointer to a `ydb_uchar_t`, the starting address of a string

Under normal circumstances `alloc ≥ used`; however, this may not be the case when a function returns a `YDB_ERR_INVSTRLEN` error. See [YDB\\_ERR\\_INVSTRLEN](#) for details.

## Macros

`YDB_ALLOC_STRING(x, strlit)` — Allocate a `ydb_string_t` structure and initialize it to `strlit`, returning the address of the structure. Note that while the `used` field is the size of `strlit`, the `alloc` field may be rounded up to a larger value.<sup>5</sup>

`YDB_FREE_STRING(x)` — Free the `ydb_string_t` structure pointed to by `x`.

## Simple API

As YottaDB local and global variables can have variable numbers of subscripts, to allow the libyottadb Simple API functions to have variable numbers of parameters, the last parameter must always be `NULL` (the standard C symbolic constant). In the definitions of functions, `ydb_string_t *varname` refers to the name of a variable, `[ydb_string_t *subscript, ...]` refers to optional subscripts following a variable name, and `NULL`); always terminates a function with optional subscripts.

Function names specific to the libyottadb Simple API end in `_s`. Others are common to both Simple API as well as the Comprehensive API.

```
ydb_status_t ydb_data_s(ydb_uint_t *value,
    ydb_string_t *varname,
    [ydb_string_t *subscript, ...]
    NULL);
```

In the location pointed to by `value`, `ydb_data_s()` returns the following information about the local or global variable node identified by `glvn`:

- 0 — There is neither a value nor a sub-tree, i.e., it is undefined.
- 1 — There is a value, but no sub-tree
- 10 — There is no value, but there is a sub-tree.
- 11 — There are both a value and a subtree.

```
ydb_status_t ydb_get_s(ydb_string_t *value,
    ydb_string_t *varname,
    [ ydb_string_t *subscript, ... ]
    NULL);
```

In the location pointed to by `value`, `ydb_get_s()` reports the value of the value of the data at the specified node.

If there is no value at the requested global or local variable node, or if the intrinsic special variable does not exist, a non-zero return value of YDB\_ERR\_GVUNDEF, YDB\_ERR\_INVSVN, or YDB\_ERR\_UNDEF indicates the error.

In a database application, a global variable node can potentially be changed by another process between the time that a process calls `ydb_length()` to get the length of the data in a node and a `ydb_get()` call to get that data. If a caller cannot ensure from the application schema that the size of the buffer it provides is large enough for a string returned by `ydb_get()`, it should code in anticipation of a potential YDB\_ERR\_INVSTRLEN return code from `ydb_get()`. See also the discussion at [YDB\\_ERR\\_INVSTRLEN](#) describing the contents of `*value` when `ydb_get-s()` returns a YDB\_ERR\_INVSTRLEN return code.

```
ydb_status_t ydb_length_s(ydb_strlen_t *value,
    ydb_string_t *varname,
    [ ydb_string_t *subscript, ... ]
    NULL);
```

In the location pointed to by `*value`, `ydb_length_s()` reports the length of the data in bytes. If the data is numeric, `*value` has the length of the canonical string representation of that value.

If there is no value at the requested global or local variable node, or if the intrinsic special variable does not exist, a non-zero return value of YDB\_ERR\_GVUNDEF, YDB\_ERR\_INVSVN, or YDB\_ERR\_UNDEF indicates the error.

```
ydb_status_t ydb_node_next_s(ydb_string_t *value,
    ydb_string_t *varname,
    [ ydb_string_t *subscript, ... ]
    NULL);
```

**In the `ydb_string_t` returns the next node in the tree in depth first search order, if one exists:**

## STUFF BELOW IS FROM OLD DOCUMENT

### Query

```
ydb_status_t ydb_alias_handle( ydb_string_t *value, ydb_varsub_t *lvn )
```

In the location pointed to by `value->address` returns the handle of the local variable referenced by `lvsub`. It is not meaningful for a caller to perform any operations on handles except to compare two handles for equality.

```
ydb_status_t ydb_data( ydb_uint_t *value, ydb_varsub_t *glvn )
```

In the location pointed to by `value`, returns the following information about the local or global variable node identified by `glvn`:

- 0 — There is neither a value nor a sub-tree, i.e., it is undefined.
- 1 — There is a value, but no sub-tree
- 10 — There is no value, but there is a sub-tree.
- 11 — There are both a value and a subtree.

The following values are only meaningful if glvn identifies a local variable node:

- 100 — The node is an alias, but there is neither a value nor a sub-tree.
- 101 — The node is an alias with a value but sub-tree.
- 110 — The node is an alias with no value, but with a sub-tree.
- 111 — The node is an alias with a value and a sub-tree.

`ydb_get( ydb_value_t *value, ydb_varsub_t *glivn )`

In the container pointer to by `value`, libyottadb returns the value referred to by `glivn`. If there is no value at the requested global or local variable node, or if the intrinsic special variable does not exist, a non-zero return value of `YDB_ERR_GVUNDEF`, `YDB_ERR_INVSVN`, or `YDB_ERR_UNDEF` indicates the error.

In a database application, a global variable node can potentially be changed by another process between the time that a process calls `ydb_length()` to get the length of the data in a node and a `ydb_get()` call to get that data. If a caller cannot ensure from the application schema that the size of the buffer it provides is large enough for a string returned by `ydb_get()`, it should code in anticipation of a potential `YDB_ERR_INVSTRLEN` return code from `ydb_get()`.

`ydb_length(ydb_ulong_t *value1, ydb_ulong_t *value2, ydb_ulong_t *value3, ydb_varsub_t *gli vn)`

For each non-null `value*` parameter, in the memory location pointed to by that parameter, libyottadb returns the following information about the node specified by `*glivsub`. If there is no value at the requested global or local variable node, or if a requested intrinsic special variable does not exist, a non-zero return value of `YDB_ERR_GVUNDEF`, `YDB_ERR_INVSVN`, or `YDB_ERR_UNDEF` indicates the error.

- `*value1` — the length of the data in bytes. If the data is numeric, `*value1` has the length of the canonical string representation of that value.
- `*value2` — if the environment variable `ydb_chset` at process startup has the (case insensitive) value “UTF-8”, `*value2` has the length of `*glivsub` in bytes; otherwise `*value2` is the same as `*value1`.
- `*value3` — if the environment variable `ydb_chset` at process startup has the (case insensitive) value “UTF-8”, `*value3` has the length of `*glivsub` in glyphs; otherwise `*value3` is the same as `*value1`.

`ydb_node_next(ydb_varsub_t *next, ydb_value_t *value, ydb_varsub_t *glvn)`

`ydb_node_next()` returns the next node in the tree in depth first search order, if one exists:

- If `next->varname->name->alloc ≥ glvn->varname->name->used` `ydb_node_next()` copies the `name->address` and `name->used` sub-fields from the `glvn->varname` structure to the `next->varname` structure, returning a `YDB_ERR_INVSTRLEN` error if `next->varname->name->alloc < glvn->varname->name->used`.
- If `next->varsub_alloc` is large enough to hold the subscripts `ydb_node_next()` sets `next->varsub_used` to the actual number of subscripts, and

## Update

## Transaction Processing

## Programming Notes

### *Dynamic typing with automatic conversion*

The YottaDB engine internally automatically converts values between numbers and strings as needed. Thus it is legitimate to lexically compare the numbers 2 and 11, with the expected result that 11 precedes 2, and it is equally legitimate to numerically compare the strings “2” and “11”, with the expected result that 11 is greater than 2 — the functions for numeric and lexical comparisons are different.

In the `ydb_value_t` structure, a caller specifies what conversion, if any, it wishes the called `libyottadb` function to perform on the return value:

- When a value is numeric, and the requested type is a string (the `tag` field is `YDB_STRING_STAR`), `libyottadb` returns the number as a canonical string in the `ydb_string_t` structure pointed to by `string_star`
- When the value is a string, and the requested type is numeric, `libyottadb` converts

When returning a string, `libyottadb` functions *always* check that the `alloc` field of the `ydb_string_t` structure is large enough for the result, returning a `YDB_ERR_STRLEN` error if it is not.

### *Numeric Considerations*

To ensure the accuracy of financial calculations, YottaDB internally stores numbers as, and performs arithmetic using, a scaled packed decimal representation with 18 significant decimal digits, with optimizations for values within a certain subset of its full range. Consequently:

- Any number that is exactly represented in YottaDB can be exactly represented as a string, with reasonably efficient conversion back and forth.
- Any integer value of up to 18 significant digits can be exactly represented by an integer type such as `ydb_longlong_t`, and integers in the inclusive range  $\pm 999,999$  are handled more efficiently than larger integers.
- In YottaDB there are numbers which can be exactly represented (such as 0.1), but which cannot be exactly represented in binary floating point.
- In 64 bit integers and binary floating point formats, there are numbers which can be exactly represented, but which cannot be exactly represented in YottaDB.

This means that for numeric keys which are not guaranteed to be integers:

- In theory, there are edge cases where a value (which would internally be in YottaDB format) returned by a function such as `ydb_subscript_next()` and converted to a `ydb_double_t` when passed back to C application code, and then converted back to YottaDB internal format in a call to `ydb_get()` can result in the node not being found because the double conversion produces a number not identical to the original. Furthermore, there is a cost to the conversion.
- Passing keys back and forth as strings avoids those edge cases, but of course still has a conversion cost.

To preserve accuracy of numeric values that are returned by `libyottadb`, and which an application code intends to simply pass back to `libyottadb` as a `libyottadb` provides a `ydb_numeric_t` type. A

value obtained from libyottadb in `ydb_numeric_t` loses no precision when returned to libyottadb, and as a further benefit is very efficient. While the actual value of `ydb_numeric_t` is opaque to application code, the `ydb_convert()` function is available.

Conversely, when passed a string that is a **canonical number** for use as a key, libyottadb automatically converts it to a number. This automatic internal conversion is irrelevant for the majority of typical application that:

- simply store and retrieve data associated with keys, potentially testing for the existence of nodes; or
- transfer keys which are numeric values between application code and libyottadb using numeric types and expect numeric ordering.

However, this automatic internal conversion does affect applications that:

- use numeric keys and expect the keys to be sorted in lexical order rather than numeric order; or
- transfer keys which are numeric values between application code and libyottadb as strings that may or may not be canonical numbers.

Applications that are affected by automatic internal conversion should prefix their keys with a character such as “x” which ensures that keys are not canonical numbers.

## Canonical Numbers

Conceptually, a canonical number is a string from the Latin character set that represents a decimal number in a standard, concise, form.

1. Any string of decimal digits, optionally preceded by a minus sign (“-“), the first of which is not “0” (except for the number zero itself), that represents an integer of no more than 18 significant digits.
  - The following are canonical numbers: “-1”, “0”, “3”, “10”, “999999999999999999”, “9999999999999999990”. Note that the last string has only 18 significant digits even though it is 19 characters long.
  - The following are not canonical numbers: “+1” (starts with “+”), “00” (has an extra leading zero), “9999999999999999999” (19 significant digits).
2. Any string of decimal digits, optionally preceded by a minus sign that includes one decimal point (“.”), the first and last of which are not “0”, that represents a number of no more than 18 significant digits.
  - The following are canonical numbers: “-.1”, “.3”, “.999999999999999999”.
  - The following are not canonical numbers “+.1” (starts with “+”), “0.3” (first digit is “0”), “.9999999999999999990” (last digit is “0”), “.9999999999999999999” (more than 18 significant digits).
3. Any of the above two forms followed by “E” followed by a canonical number integer in the range -43 to +47 such that the magnitude of the resulting number is between 1E-43 through 1E47.

## Tokens



Since numeric and non-numeric subscripts can be freely intermixed in YottaDB, it requires knowledge of the application schema to know whether an application mixes numeric and string subscripts at the same level for a variable.

*Consider whether this can be deferred for an initial implementation.*

- 
- 1 Note for implementers: the actual values are negated ZMESSAGE error codes.
  - 2 Note for implementers: YottaDB silently truncates local variable names that are too long. The implementation should catch this and return an error code, e.g., something like `YDB_ERR_LVINVALID`.
  - 3 Note for implementers: under the covers, this is `UNDEF` but renamed to be more meaningful.
  - 4 Strings in YottaDB are arbitrary sequences of bytes that are not null-terminated. Other languages may refer to them as binary data or blobs.
  - 5 Note for implementers: under the covers, `YDB_ALLOC_STRING()` and `YDB_FREE_STRING()` should call the `ydb_malloc()` and `ydb_free()` functions, which are aliases for the `gtm_malloc()` and `gtm_free()` functions (i.e., either prefix calls the same function).